



**UNIVERSIDAD  
NACIONAL  
DE INGENIERÍA**

**Facultad de Ingeniería  
Industrial y de  
Sistemas**

**SECC. A**

# **CURSO REDES NEURONALES Y APRENDIZAJE PROFUNDO**

## **TAREA 1**

### **INTEGRANTES GRUPO 8**

Kevin Gómez Villanueva  
Fernando Boza Gutarra


Umbert Lewis de la Cruz Rodriguez  
Yovany Romero Ramos

1. Hacer las operaciones de forward y backward propagation de forma manual para un MLP que tome 3 entradas (3x1), tenga 2 hidden layers de tamaño 4, y una salida de 1x1. El cálculo lo deben hacer para la siguiente data:

$$X_s = \begin{bmatrix} 2.5 & 3.5 & -0.5 \\ 4.0 & -1.0 & 0.5 \\ 0.5 & 1.5 & 1.0 \\ 3.0 & 2.0 & -1.5 \end{bmatrix}$$

$$y_s = \begin{bmatrix} 1.0 \\ -1.0 \\ -1.0 \\ 1.0 \end{bmatrix}$$

# Modelamiento de Red Neuronal



Capa de entrada

Capas ocultas

Capa de salida

# Proceso ForWard

$$z_1 = W_1 \cdot X + b_1 \quad a_1 = \tanh(z_1)$$

$$z_2 = W_2 \cdot a_1 + b_2 \quad a_2 = \tanh(z_2)$$

$$z_3 = W_3 \cdot a_2 + b_3 \quad a_3 = \tanh(z_3) = \hat{y}$$

# Pesos y Bias

$$W_1 = \begin{bmatrix} -0.341 & 0.488 & -0.579 \\ -0.741 & -0.542 & -0.716 \\ -0.432 & -0.389 & -0.726 \\ 0.854 & -0.244 & -0.297 \end{bmatrix}$$

$$b_1 = \begin{bmatrix} 0.421 \\ -0.068 \\ 0.392 \\ 0.079 \end{bmatrix}$$

$$W_2 = \begin{bmatrix} 0.824 & -0.793 & 0.857 & 0.389 \\ -0.840 & 0.601 & -0.711 & -0.538 \\ -0.962 & -0.148 & 0.315 & 0.844 \\ 0.759 & -0.675 & -0.054 & 0.430 \end{bmatrix}$$

$$b_2 = \begin{bmatrix} 0.281 \\ 0.671 \\ 0.275 \\ 0.209 \end{bmatrix}$$

$$W_3 = [-0.355 \quad -0.301 \quad 0.569 \quad -0.967]$$

$$b_3 = [0.998]$$

# Resultados Forward

<b>z1</b>	[[1.564, -1.721, 0.403, 1.241], [-3.459, -2.847, -1.967, -2.300], [-1.685, -1.309, -1.133, -0.591], [1.508, 3.592, -0.158, 2.599]]
<b>a1</b>	[[0.916, -0.938, 0.382, 0.846], [-0.998, -0.993, -0.962, -0.980], [-0.934, -0.864, -0.812, -0.531], [0.907, 0.998, -0.157, 0.989]]
<b>z2</b>	[[1.380, -0.056, 0.601, 1.685], [-0.522, 0.939, 0.433, -0.783], [0.013, 1.895, -0.338, 0.274], [2.019, 0.644, 1.125, 1.967]]
<b>a2</b>	[[0.881, -0.056, 0.538, 0.933], [-0.479, 0.735, 0.408, -0.654], [0.013, 0.956, -0.326, 0.268], [0.965, 0.568, 0.809, 0.962]]
<b>z3</b>	[[-0.096, 0.792, -0.284, 0.087]]
<b>a3 = ypred</b>	[[-0.096, 0.659, -0.276, 0.087]]

# Proceso BackWard

$$\text{Loss} = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

$$\text{Error} = \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

$$w' = w - \eta \frac{\partial \text{Error}}{\partial w}$$

$$b' = b - \eta \sum_{i=1}^n \frac{\partial \text{Error}}{\partial b}$$

# Proceso BackWard

$$\delta_3 = 2.(\hat{y} - y). (1 - \tanh^2(z_3))$$

$$\delta_2 = W_3. \delta_3. (1 - \tanh^2(z_2))$$

$$\delta_1 = W_2. \delta_2. (1 - \tanh^2(z_1))$$

$$W'_1 = W_1 - 0.1\delta_1. X_s$$

$$W'_2 = W_2 - 0.1\delta_2. a_1$$

$$W'_3 = W_3 - 0.1\delta_3. a_2$$

$$b'_1 = b_1 - 0.1 \sum \delta_1$$

$$b'_2 = b_2 - 0.1 \sum \delta_2$$

$$b'_3 = b_3 - 0.1 \sum \delta_3$$

# Resultados BackWard

W'1	[[ -0.34554259 0.4935403 -0.44262008] [ -0.7502685 -0.54508613 -0.71749508] [ -0.28692567 -0.3123435 -0.78227928] [ 0.8926412 -0.22996956 -0.35290302]]
W'2	[[ 0.75195582 -0.86552398 0.79253729 0.42652176] [ -0.92420782 0.62329247 -0.69687519 -0.59349348] [ -0.78492013 -0.29102007 0.2121559 1.05201748] [ 0.63649636 -0.81285838 -0.17622029 0.51966873]]
W'3	[[ -0.05558177 -0.71627873 0.48484404 -0.79736761]]
b'1	[[ 0.48137072] [ -0.07279767] [ 0.43346154] [ 0.05973635]]
b'2	[[ 0.3552724 ] [ 0.64854159] [ 0.41713561] [ 0.34912214]]
b'3	[[ 1.07507519]]



# Código

```
[ ] # Capa Oculta 1
z1 = W1.dot(X.T) + b1
a1 = tanh(z1)
print("z1:", z1)
print("a1:", a1)

z1: [[ 1.56370189 -1.72053842  0.40277158  1.24051351]
 [-3.45886626 -2.84680452 -1.96743566 -2.30046325]
 [-1.68502365 -1.3086783  -1.13335609 -0.59149301]
 [ 1.50819259  3.59206435 -0.15816039  2.59942234]]
a1: [[ 0.91601815 -0.93792783  0.38231794  0.84560199]
 [-0.99802181 -0.99328777 -0.96165318 -0.98011465]
 [-0.93351042 -0.8639406  -0.81216475 -0.53096855]
 [ 0.90661778  0.99848409 -0.15685467  0.98901479]]

# Capa Oculta 2
z2 = W2.dot(a1) + b2
a2 = tanh(z2)
print("z2:", z2)
print("a2:", a2)

z2: [[ 1.37989756 -0.0562545  0.60110234  1.68462892]
 [-0.52225653  0.93933922  0.43348545 -0.78275412]
 [ 0.01302297  1.89478974 -0.33821991  0.2744689 ]
 [ 2.0186989  0.64418511  1.1252113  1.96676804]]
a2: [[ 0.88092833 -0.05619523  0.5378335  0.93345966]
 [-0.47943973  0.73491851  0.40823003 -0.65428466]
 [ 0.01302223  0.95578918 -0.32588728  0.26777823]
 [ 0.96532513  0.56774241  0.80937396  0.96160293]]

# Capa de Salida
z3 = W3.dot(a2) + b3
a3 = tanh(z3)
print("z3:", z3)
print("a3:", a3)

z3: [[-0.09580384  0.79180339 -0.28356452  0.08680918]]
a3: [[-0.0955118  0.65942944 -0.2762009  0.08659178]]
```

```
[ ] error = a3 - y.T
delta3 = 2 * error * tanh_derivative(z3)
print("error:", error)
print("delta3:", delta3)

error: [[-1.0955118  1.65942944  0.7237991 -0.91340822]]
delta3: [[-2.17103598  1.87566244  1.33716536 -1.81311873]]

loss = np.sum((y.T - a3)**2)
print(loss)

5.312051880293905

delta2 = W3.T.dot(delta3) * tanh_derivative(z2)
print("delta2:", delta2)

delta2: [[ 0.17247053 -0.66320567 -0.33710103  0.08273963]
 [ 0.50348499 -0.25975528 -0.33555434  0.31225281]
 [-1.2354633  0.09230852  0.68023774 -0.95796281]
 [ 0.14300552 -1.2285945 -0.44579195  0.13199941]]

delta1 = W2.T.dot(delta2) * tanh_derivative(z1)
print("delta1:", delta1)

delta1: [[ 0.16346314 -0.16233431 -0.84395226  0.23575742]
 [ 0.00099784  0.01586053  0.01998641  0.00689013]
 [-0.07805895 -0.07299838  0.06407392 -0.33038147]
 [-0.21089959 -0.00172385  0.42081974 -0.01938593]]
```

```
[ ] learning_rate = 0.1
W1 -= learning_rate * dW1
b1 -= learning_rate * db1
W2 -= learning_rate * dW2
b2 -= learning_rate * db2
W3 -= learning_rate * dW3
b3 -= learning_rate * db3

print("\nUpdated Weights and Biases:")
print("W1: {W1}")
print("W2: {W2}")
print("W3: {W3}")
print("b1: {b1}")
print("b2: {b2}")
print("b3: {b3}")

Updated Weights and Biases:
W1: [[-0.34554259  0.4935403 -0.44262008]
 [-0.7502685 -0.54508613 -0.71749508]
 [-0.28692567 -0.3123435 -0.78227928]
 [ 0.8926412 -0.22996956 -0.35290302]]
W2: [[ 0.75195582 -0.86552398  0.79253729  0.42652176]
 [-0.92420782  0.62329247 -0.69687519 -0.59349348]
 [-0.78492013 -0.29102007  0.2121559  1.05201748]
 [ 0.63649636 -0.81285838 -0.17622029  0.51966873]]
W3: [[-0.05558177 -0.71627873  0.48484404 -0.79736761]]
b1: [[ 0.48137072]
 [-0.07279767]
 [ 0.43346154]
 [ 0.05973635]]
b2: [[0.3552724 ]
 [0.64854159]
 [0.41713561]
 [0.34912214]]
b3: [[1.07507519]]
```

## 2. Verificar el resultado del cálculo a mano usando la librería vista en la segunda clase (Micrograd). Hacer lo mismo usando PyTorch.

**Micrograd** es una librería que define una clase `Value` que se usa para representar valores numéricos en una red de cómputo que admite operaciones como suma, multiplicación y otras funciones, junto con la capacidad de calcular gradientes para el aprendizaje automático mediante la diferenciación automática.

```
class Value:

    def __init__(self, data, _children=(), _op='', label=''):
        self.data = data
        self.grad = 0.0
        self._backward = lambda: None
        self._prev = set(_children)
        self._op = _op
        self.label = label
```

# Otras clases usadas

```
import random
class Neuron:

    def __init__(self, nin):
        self.w = [Value(random.uniform(-1,1)) for _ in range(nin)]
        print("W:", self.w)
        self.b = Value(random.uniform(-1,1))
        print("b:", self.b)

    def __call__(self, x):
        # w * x + b
        act = sum((wi*xi for wi, xi in zip(self.w, x)), self.b)
        out = act.tanh()
        return out

    def parameters(self):
        return self.w + [self.b]

class Layer:

    def __init__(self, nin, nout):
        self.neurons = [Neuron(nin) for _ in range(nout)]

    def __call__(self, x):
        outs = [n(x) for n in self.neurons]
        return outs[0] if len(outs) == 1 else outs

    def parameters(self):
        return [p for neuron in self.neurons for p in neuron.parameters()]
```

```
class MLP:

    def __init__(self, nin, nouts):
        sz = [nin] + nouts
        self.layers = [Layer(sz[i], sz[i+1]) for i in range(len(nouts))]

    def __call__(self, x):
        for layer in self.layers:
            x = layer(x)
        return x

    def parameters(self):
        return [p for layer in self.layers for p in layer.parameters()]
```

- La clase **Neuron** representa una sola neurona en una red neuronal.
- La clase **Layer** representa una capa de varias neuronas.
- La clase **MLP (Multi-Layer Perceptron)** representa una red neuronal compuesta por varias capas.



# Red Neuronal MultiCapa

```
[ ] n = MLP(3, [4, 4, 1])
```

```
→ W: [Value(data=-0.34108091502837, grad=0.0), Value(data=0.4875444774924518, grad=0.0), Value(data=-0.5786687891714137, grad=0.0)]  
b: Value(data=0.42066411352123567, grad=0.0)  
W: [Value(data=-0.7406084700295616, grad=0.0), Value(data=-0.5419469493214886, grad=0.0), Value(data=-0.7157868250555495, grad=0.0)]  
b: Value(data=-0.06842417572242421, grad=0.0)  
W: [Value(data=-0.4315505068114813, grad=0.0), Value(data=-0.3888294993216346, grad=0.0), Value(data=-0.7260616380268912, grad=0.0)]  
b: Value(data=0.3917250492343052, grad=0.0)  
W: [Value(data=0.8544519713956533, grad=0.0), Value(data=-0.24436625968891335, grad=0.0), Value(data=-0.29745436944045234, grad=0.0)]  
b: Value(data=0.07861738679252617, grad=0.0)  
W: [Value(data=0.8240668394757076, grad=0.0), Value(data=-0.7925535106179675, grad=0.0), Value(data=0.856718957818575, grad=0.0), Value(data=0.38940887530348656, grad=0.0)]  
b: Value(data=0.2807627460701372, grad=0.0)  
W: [Value(data=-0.8401491947426882, grad=0.0), Value(data=0.6005090770290116, grad=0.0), Value(data=-0.710761829952594, grad=0.0), Value(data=-0.5376371931922113, grad=0.0)]  
b: Value(data=0.6705844066805486, grad=0.0)  
W: [Value(data=-0.9617475006414558, grad=0.0), Value(data=-0.1484109682541901, grad=0.0), Value(data=0.315131084013728, grad=0.0), Value(data=0.843811249537862, grad=0.0)]  
b: Value(data=0.2750476238785533, grad=0.0)  
W: [Value(data=0.7589476929906058, grad=0.0), Value(data=-0.6751635852663345, grad=0.0), Value(data=-0.05422983414165872, grad=0.0), Value(data=0.430008053438502, grad=0.0)]  
b: Value(data=0.20918398742213418, grad=0.0)  
W: [Value(data=-0.3547048938594579, grad=0.0), Value(data=-0.3011280508308256, grad=0.0), Value(data=0.5691627652457749, grad=0.0), Value(data=-0.966577207820160, grad=0.0)]  
b: Value(data=0.9979425023631463, grad=0.0)
```

# Resultados

Iteración 0

Predicciones: [-0.09551180109351012, 0.6594294380423391, -0.2762009023561765, 0.08659177780839312]

Pesos y bias actualizados:

Parámetro 0: -0.3455425864395133

Parámetro 1: 0.49354030472046384

Parámetro 2: -0.44262007762858524

Parámetro 3: 0.4813707157639071

Parámetro 4: -0.7502685003738252

Parámetro 5: -0.5450861251164729

Parámetro 6: -0.7174950822443538

Parámetro 7: -0.07279766624127787

Parámetro 8: -0.2869256709509327

Parámetro 9: -0.3123434983594058

Parámetro 10: -0.7822792789756129

Parámetro 11: 0.43346153792328573

Parámetro 12: 0.8926412032411102

Parámetro 13: -0.22996956288809678

Parámetro 14: -0.35290302011138436

Parámetro 15: 0.0597363504502513

Parámetro 16: 0.7519558171679739

Parámetro 17: -0.865523978704413

Parámetro 18: 0.7925372877736636

Parámetro 19: 0.4265217627334761

Parámetro 20: 0.3552723998299817

Parámetro 21: -0.9242078198017564

Parámetro 22: 0.623292469237169

Parámetro 23: -0.696875192571201

Parámetro 24: -0.593493477673018

Parámetro 25: 0.6485415878541207

Parámetro 26: -0.7849201299364432

Parámetro 27: -0.2910200672775243

Parámetro 28: 0.21215590477957946

Parámetro 29: 1.0520174753370677

Parámetro 30: 0.4171356087187202

Parámetro 31: 0.6364963596437079

Parámetro 32: -0.8128583812158143

Parámetro 33: -0.17622028514159255

Parámetro 34: 0.5196687337910336

Parámetro 35: 0.3491221402067136

Parámetro 36: -0.055581770552016574

Parámetro 37: -0.7162787275772213

Parámetro 38: 0.48484404322738217

Parámetro 39: -0.797367614376691

Parámetro 40: 1.075075192880589

La clase MLP define una red neuronal usando `torch.nn.Module`, que es la base para construir modelos en PyTorch.

- **La primera capa** toma una entrada de tamaño 3 y produce una salida de tamaño 4.
- **La segunda capa** toma una entrada de tamaño 4 y produce otra salida de tamaño 4.
- **La capa de salida** toma una entrada de tamaño 4 y produce una única salida.

```
class MLP(nn.Module):
    def __init__(self, W1, b1, W2, b2, W3, b3):
        super(MLP, self).__init__()
        self.layer1 = nn.Linear(3, 4)
        self.layer2 = nn.Linear(4, 4)
        self.output_layer = nn.Linear(4, 1)

        with torch.no_grad():
            self.layer1.weight = nn.Parameter(W1)
            self.layer1.bias = nn.Parameter(b1)
            self.layer2.weight = nn.Parameter(W2)
            self.layer2.bias = nn.Parameter(b2)
            self.output_layer.weight = nn.Parameter(W3)
            self.output_layer.bias = nn.Parameter(b3)

    def forward(self, x):
        x = torch.tanh(self.layer1(x))
        x = torch.tanh(self.layer2(x))
        x = torch.tanh(self.output_layer(x))
        return x
```

# Resultados

```
print("\nPesos y bias después de la retropropagación:")
for name, param in model.named_parameters():
    print(f"{name}: {param.data}")

print(f"\nPérdida final: {loss}")
```

```
Pesos y bias después de la retropropagación:
layer1.weight: tensor([[ -0.3652,  0.4924, -0.5307],
 [ -0.7427, -0.5432, -0.7167],
 [ -0.4137, -0.3822, -0.7352],
 [  0.8516, -0.2657, -0.3195]], device='cuda:0')
layer1.bias: tensor([[ 0.4201],
 [ -0.0663],
 [  0.3961],
 [  0.0718]], device='cuda:0')
layer2.weight: tensor([[ 0.8030, -0.8150,  0.8352,  0.4034],
 [ -0.8521,  0.5926, -0.7187, -0.5403],
 [ -0.9284, -0.1723,  0.2961,  0.8872],
 [  0.7380, -0.7009, -0.0788,  0.4464]], device='cuda:0')
layer2.bias: tensor([[0.2956],
 [0.7166],
 [0.2862],
 [0.2188]], device='cuda:0')
output_layer.weight: tensor([[ -0.3148, -0.4109,  0.5292, -0.9650]], device='cuda:0')
output_layer.bias: tensor([[0.9839]], device='cuda:0')

Pérdida final: 1.1816823482513428
```



### 3. Usando la data de Kaggle sobre Lung Cancer ( Histopathological Images), clasificar dado la imagen está en uno de estas clases: ['adenocarcinoma', 'benign', 'squamous cell carcinoma'].

#### Lung and Colon Cancer Histopathological Image Dataset (LC25000)

Andrew A. Borkowski, MD<sup>\*1,2</sup>, Marilyn M. Bui, MD, PhD<sup>2,3</sup>, L. Brannon Thomas, MD, PhD<sup>1,2</sup>,  
Catherine P. Wilson, MT<sup>1</sup>, Lauren A. DeLand, RN<sup>1</sup>, Stephen M. Mastorides, MD<sup>1,2</sup>

El conjunto de datos contiene 15000 imágenes en color, divididas en 3 clases de 5000 imágenes cada una. Todas las imágenes tienen un tamaño de 768 x 768 píxeles y están en formato jpeg. Las clases son adenocarcinomas de pulmón, carcinomas de células escamosas de pulmón y tejidos pulmonares benignos.



# Lung and Colon Cancer Histopathological Images

246


New Notebook

Download (2 GB)




Data Card Code (196) Discussion (0) Suggestions (2)


colon\_aca (5000 files)




colonca1.jpeg  
118.75 kB




colonca10.jpeg  
77.29 kB




colonca100.jpeg  
87.18 kB




colonca1000.jpeg  
111.94 kB




colonca1001.jpeg  
94.43 kB




colonca1002.jpeg  
84.72 kB




colonca1003.jpeg  
102.05 kB




colonca1004.jpeg  
82.94 kB




colonca1005.jpeg  
80.19 kB




colonca1006.jpeg  
108.65 kB




colonca1007.jpeg  
79.33 kB




colonca1008.jpeg  
83.01 kB



colonca1009.jpeg  
100.93 kB



colonca101.jpeg  
118.55 kB



colonca1010.jpeg  
85.24 kB

Data Explorer


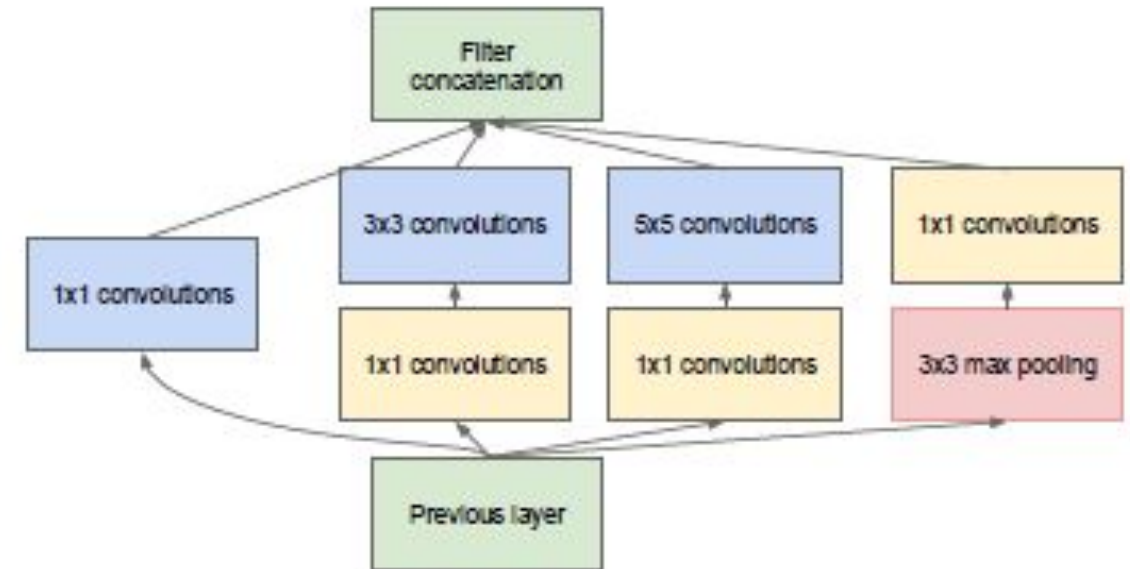
Version 1 (1.89 GB)

- lung\_colon\_image\_set
  - colon\_image\_sets
    - colon\_aca
      - colonca1.jpeg
      - colonca10.jpeg
      - colonca100.jpeg
      - colonca1000.jpeg
      - colonca1001.jpeg
      - colonca1002.jpeg
      - colonca1003.jpeg
      - colonca1004.jpeg
      - colonca1005.jpeg
      - colonca1006.jpeg
      - colonca1007.jpeg
      - colonca1008.jpeg
      - colonca1009.jpeg
      - colonca101.jpeg
      - colonca1010.jpeg
      - colonca1011.jpeg
      - colonca1012.jpeg
      - colonca1013.jpeg

# Inceptionv1

## Módulo Inception

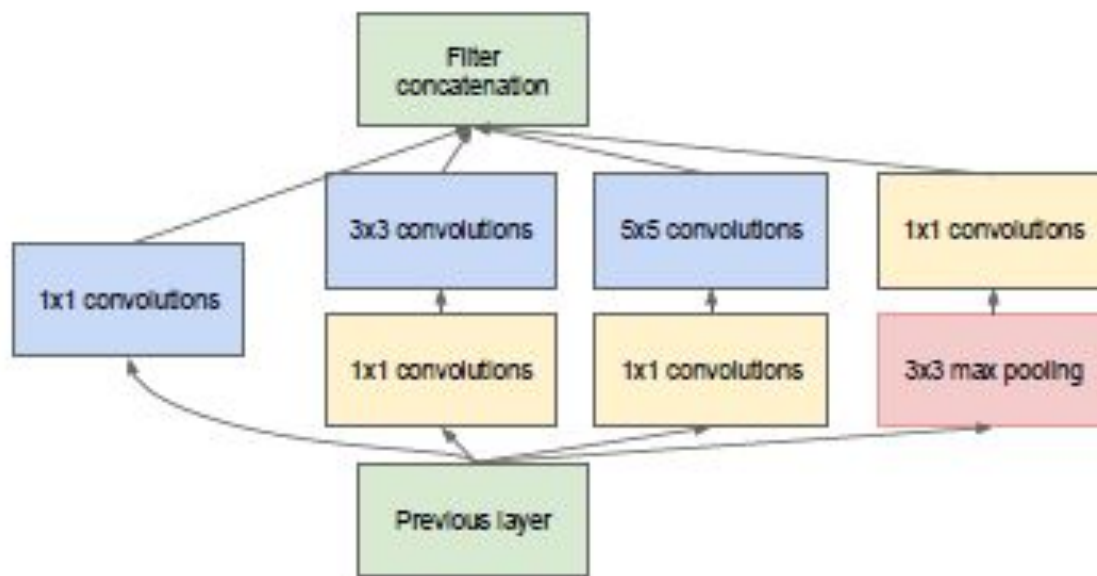
- convolution
- max pooling
- channel concatenation
- channel-wise normalization
- fully-connected layer
- softmax





# Inceptionv1

## Módulo Inception



```
class Inception_block(nn.Module):
    def __init__(self, in_channels, out_1x1, red_3x3, out_3x3, red_5x5, out_5x5, out_1x1pool):
        super(Inception_block, self).__init__()

        self.branch1 = conv_block(in_channels, out_1x1, kernel_size=1)

        self.branch2 = nn.Sequential(
            conv_block(in_channels, red_3x3, kernel_size=1),
            conv_block(red_3x3, out_3x3, kernel_size=3, stride=1, padding=1)
        )

        self.branch3 = nn.Sequential(
            conv_block(in_channels, red_5x5, kernel_size=1),
            conv_block(red_5x5, out_5x5, kernel_size=5, stride=1, padding=2)
        )

        self.branch4 = nn.Sequential(
            nn.MaxPool2d(kernel_size=3, stride=1, padding=1),
            conv_block(in_channels, out_1x1pool, kernel_size=1)
        )

    def forward(self, x):
        return torch.cat([self.branch1(x), self.branch2(x), self.branch3(x), self.branch4(x)], 1)
```

# Inceptionv1

```
class GoogleNet(nn.Module):
    def __init__(self, in_channels=3, num_classes=1000):
        super(GoogleNet, self).__init__()

        self.conv1 = conv_block(in_channels=in_channels, out_channels=64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3))
        self.maxpool1 = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)

        self.conv2 = conv_block(in_channels=64, out_channels=192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        self.maxpool2 = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)

        self.inception3a = Inception_block(192, 64, 96, 128, 16, 32, 32)
        self.inception3b = Inception_block(256, 128, 128, 192, 32, 96, 64)
        self.maxpool3 = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)

        self.inception4a = Inception_block(480, 192, 96, 208, 16, 48, 64)
        self.inception4b = Inception_block(512, 160, 112, 224, 24, 64, 64)
        self.inception4c = Inception_block(512, 128, 128, 256, 24, 64, 64)
        self.inception4d = Inception_block(512, 112, 144, 288, 32, 64, 64)
        self.inception4e = Inception_block(528, 256, 160, 320, 32, 128, 128)
        self.maxpool4 = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)

        self.inception5a = Inception_block(832, 256, 160, 320, 32, 128, 128)
        self.inception5b = Inception_block(832, 384, 192, 384, 48, 128, 128)

        self.avpool = nn.AvgPool2d(kernel_size=16)
        self.dropout = nn.Dropout(p=0.4)
        self.fc1 = nn.Linear(1024, num_classes)
```

```
def forward(self, x):
    x = self.conv1(x)
    x = self.maxpool1(x)
    x = self.conv2(x)
    x = self.maxpool2(x)

    x = self.inception3a(x)
    x = self.inception3b(x)
    x = self.maxpool3(x)

    x = self.inception4a(x)
    x = self.inception4b(x)
    x = self.inception4c(x)
    x = self.inception4d(x)
    x = self.inception4e(x)
    x = self.maxpool4(x)

    x = self.inception5a(x)
    x = self.inception5b(x)
    x = self.avpool(x)

    x = x.reshape(x.shape[0], -1)
    x = self.dropout(x)
    x = self.fc1(x)


    return x
```

# Inceptionv1

```
Epoch [1/10], Loss: 0.4093  
Validation Loss: 0.1961, Accuracy: 92.98%  
Epoch [2/10], Loss: 0.2671  
Validation Loss: 0.2879, Accuracy: 87.20%  
Epoch [3/10], Loss: 0.2284  
Validation Loss: 0.1590, Accuracy: 92.93%  
Epoch [4/10], Loss: 0.1980  
Validation Loss: 0.3810, Accuracy: 81.11%  
Epoch [5/10], Loss: 0.1699  
Validation Loss: 0.1200, Accuracy: 95.20%  
Epoch [6/10], Loss: 0.1561  
Validation Loss: 0.1104, Accuracy: 95.33%  
Epoch [7/10], Loss: 0.1504  
Validation Loss: 0.1222, Accuracy: 94.62%  
Epoch [8/10], Loss: 0.1348  
Validation Loss: 0.1032, Accuracy: 95.78%  
Epoch [9/10], Loss: 0.1220  
Validation Loss: 0.0947, Accuracy: 96.00%  
Epoch [10/10], Loss: 0.1131  
Validation Loss: 0.0743, Accuracy: 97.24%  
Evaluando el modelo en el conjunto de test...  
Validation Loss: 0.0687, Accuracy: 97.69%
```

Epoch: 10

Accuracy: 97.69%



Predicción: adenocarcinoma con 97.00% confidence

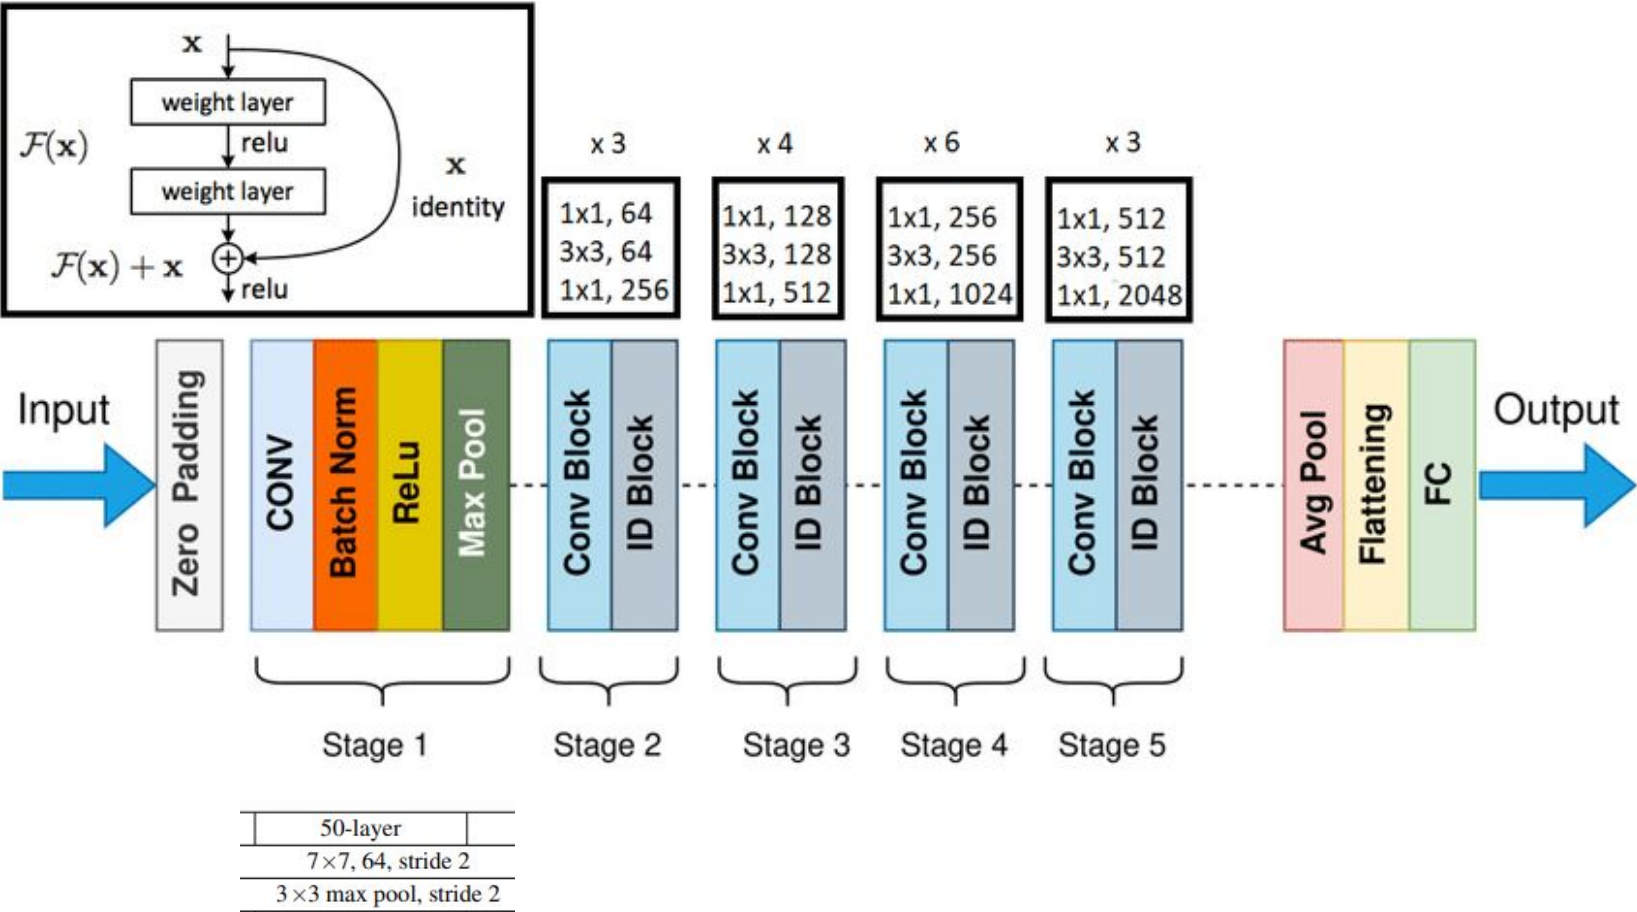


# ResNet-50

Previene el degradación de la gradiente en el backpropagation



$$H(x) = F(x) + x$$

$$\frac{\partial H(x)}{\partial x} = \frac{\partial F(x)}{\partial x} + 1$$



layer name	output size	50-layer
conv1	112×112	7×7, 64, stride 2
conv2_x	56×56	3×3 max pool, stride 2
		<div><div>1×1, 64</div><div>3×3, 64</div><div>1×1, 256</div></div> ×3
conv3_x	28×28	<div><div>1×1, 128</div><div>3×3, 128</div><div>1×1, 512</div></div> ×4
conv4_x	14×14	<div><div>1×1, 256</div><div>3×3, 256</div><div>1×1, 1024</div></div> ×6
conv5_x	7×7	<div><div>1×1, 512</div><div>3×3, 512</div><div>1×1, 2048</div></div> ×3
	1×1	average pool, 1000-d fc, s
FLOPs		3.8×10 <sup>9</sup>

# ResNet-50



50-layer
$7 \times 7$ , 64, stride 2
$3 \times 3$ max pool, stride 2


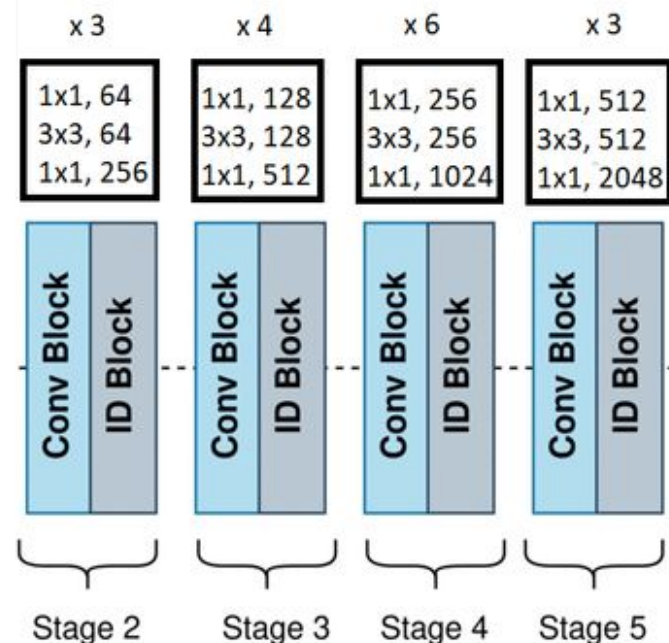
```
class ResNet50(nn.Module):
    def __init__(self, block, layers, num_classes=3):
        super(ResNet50, self).__init__()
        self.in_channels = 64

        self.conv1 = nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3, bias=False) # First layer with large 7x7 kernel
        self.bn1 = nn.BatchNorm2d(64)
        self.relu = nn.ReLU(inplace=True)
        self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)

        self.layer1 = self._make_layer(block, 64, layers[0])
        self.layer2 = self._make_layer(block, 128, layers[1], stride=2)
        self.layer3 = self._make_layer(block, 256, layers[2], stride=2)
        self.layer4 = self._make_layer(block, 512, layers[3], stride=2)

        self.avgpool = nn.AdaptiveAvgPool2d((1, 1))
        self.fc = nn.Linear(512 * 4, num_classes)
```

# ResNet-50



```
class Bottleneck(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1, downsample=None):
        super(Bottleneck, self).__init__()
        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=1, bias=False) # 1x1 convolution
        self.bn1 = nn.BatchNorm2d(out_channels)
        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3, stride=stride, padding=1, bias=False) # 3x3 convolution
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.conv3 = nn.Conv2d(out_channels, out_channels * 4, kernel_size=1, bias=False) # 1x1 convolution to increase channels
        self.bn3 = nn.BatchNorm2d(out_channels * 4)
        self.relu = nn.ReLU(inplace=True)

        self.downsample = downsample

    def forward(self, x):
        identity = x

        out = self.conv1(x)
        out = self.bn1(out)
        out = self.relu(out)

        out = self.conv2(out)
        out = self.bn2(out)
        out = self.relu(out)

        out = self.conv3(out)
        out = self.bn3(out)

        if self.downsample is not None:
            identity = self.downsample(x)

        out += identity
        out = self.relu(out)

        return out
```

```
class ResNet50(nn.Module):
    def __init__(self, block, layers, num_classes=3):
        super(ResNet50, self).__init__()
        self.in_channels = 64


        self.conv1 = nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3, bias=False)
        self.bn1 = nn.BatchNorm2d(64)
        self.relu = nn.ReLU(inplace=True)
        self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)

        self.layer1 = self._make_layer(block, 64, layers[0])
        self.layer2 = self._make_layer(block, 128, layers[1], stride=2)
        self.layer3 = self._make_layer(block, 256, layers[2], stride=2)
        self.layer4 = self._make_layer(block, 512, layers[3], stride=2)

        self.avgpool = nn.AdaptiveAvgPool2d((1, 1))
        self.fc = nn.Linear(512 * 4, num_classes)
```



# ResNet-50



```
class ResNet50(nn.Module):
    def __init__(self, block, layers, num_classes=3):
        super(ResNet50, self).__init__()
        self.in_channels = 64

        self.conv1 = nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3, bias=False) # First layer with large 7x7 kernel
        self.bn1 = nn.BatchNorm2d(64)
        self.relu = nn.ReLU(inplace=True)
        self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)

        self.layer1 = self._make_layer(block, 64, layers[0])
        self.layer2 = self._make_layer(block, 128, layers[1], stride=2)
        self.layer3 = self._make_layer(block, 256, layers[2], stride=2)
        self.layer4 = self._make_layer(block, 512, layers[3], stride=2)

        self.avgpool = nn.AdaptiveAvgPool2d((1, 1))
        self.fc = nn.Linear(512 * 4, num_classes)
```

```
def forward(self, x):
    x = self.conv1(x)
    x = self.bn1(x)
    x = self.relu(x)
    x = self.maxpool(x)

    x = self.layer1(x)
    x = self.layer2(x)
    x = self.layer3(x)
    x = self.layer4(x)


    x = self.avgpool(x)
    x = x.view(x.size(0), -1) # Flatten the output
    x = self.fc(x)

    return x
```

# ResNet-50

```
→ /usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.  
  warnings.warn(_create_warning_msg(  
Epoch [1/10], Loss: 0.3610, Accuracy: 86.28%  
Validation Accuracy: 81.07%  
Epoch [2/10], Loss: 0.2721, Accuracy: 89.44%  
Validation Accuracy: 89.51%  
Epoch [3/10], Loss: 0.2294, Accuracy: 90.74%  
Validation Accuracy: 90.89%  
Epoch [4/10], Loss: 0.1994, Accuracy: 91.89%  
Validation Accuracy: 90.93%  
Epoch [5/10], Loss: 0.2032, Accuracy: 91.87%  
Validation Accuracy: 95.69%  
Epoch [6/10], Loss: 0.1891, Accuracy: 92.97%  
Validation Accuracy: 88.40%  
Epoch [7/10], Loss: 0.2003, Accuracy: 92.85%  
Validation Accuracy: 95.42%  
Epoch [8/10], Loss: 0.1381, Accuracy: 94.67%  
Validation Accuracy: 96.89%  
Epoch [9/10], Loss: 0.1226, Accuracy: 95.29%  
Validation Accuracy: 97.42%  
Epoch [10/10], Loss: 0.1182, Accuracy: 95.98%  
Validation Accuracy: 98.00%  
Accuracy on test set: 98.00%
```

```
image_path = './data/test/adenocarcinoma/0096.jpg'  
predicted_class, confidence, probabilities = predict_image(image_path, model)  
  
classes = ['adenocarcinoma', 'benign', 'squamous_cell_carcinoma']  
  
print(f'Predicción: {classes[predicted_class]} con {confidence*100:.2f}% confidence')
```



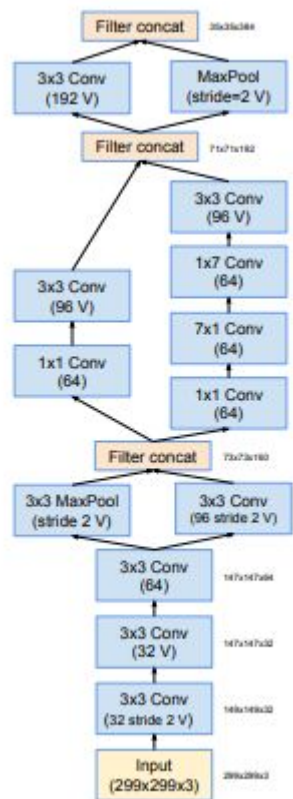
Predicción: adenocarcinoma con 98.48% confidence

Epoch: 10


Accuracy: 98%

# Inception + Resnetv2

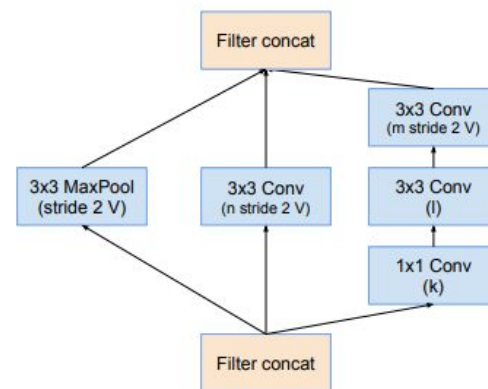
## STEM



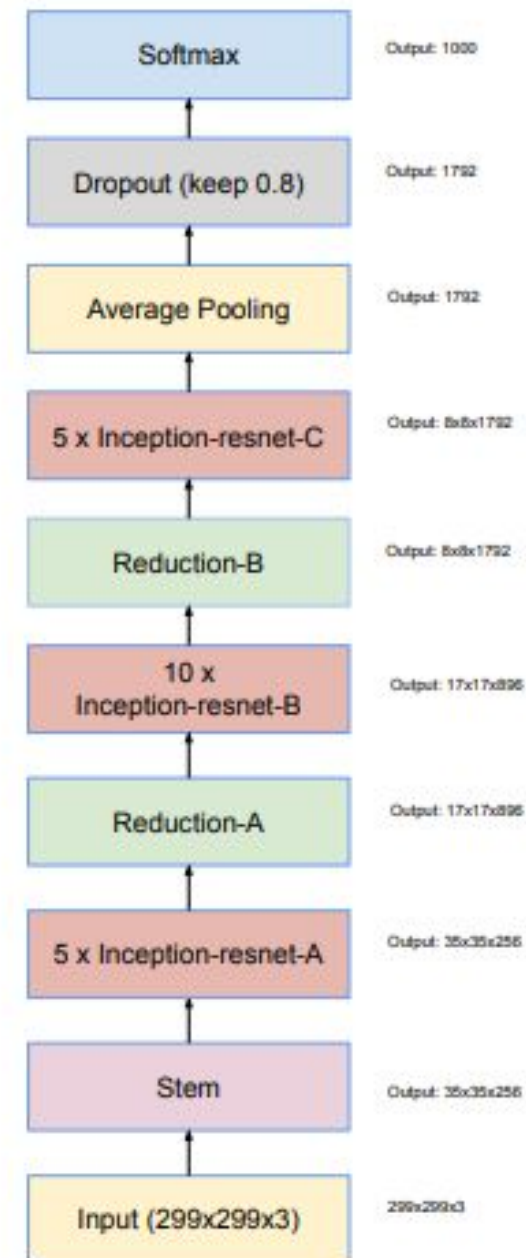

## Inception-Resnet A



## Reduction A




## Reduction B



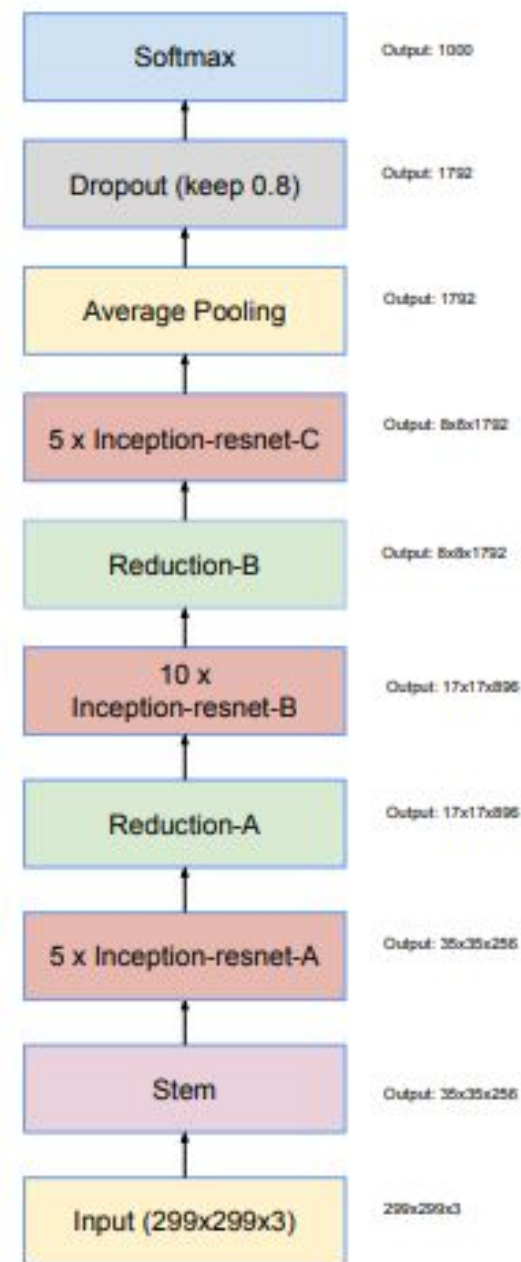
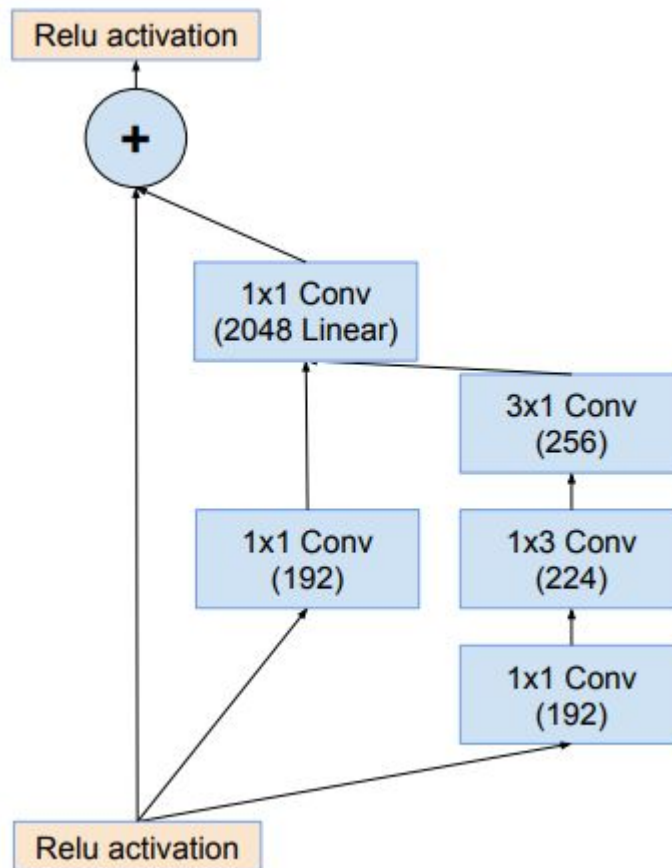


# Inception + Resnetv2

Inception-Resnet B



Inception-Resnet C



# Inception + Resnetv2

```
class Stem(nn.Module):
    def __init__(self, in_channels):
        super(Stem, self).__init__()
        self.features = nn.Sequential(
            Conv2d(in_channels, 32, 3, stride=2, padding=0, bias=False), # 149 x 149 x 32
            Conv2d(32, 32, 3, stride=1, padding=0, bias=False), # 147 x 147 x 32
            Conv2d(32, 64, 3, stride=1, padding=1, bias=False), # 147 x 147 x 64
            nn.MaxPool2d(3, stride=2, padding=0), # 73 x 73 x 64
            Conv2d(64, 80, 1, stride=1, padding=0, bias=False), # 73 x 73 x 80
            Conv2d(80, 192, 3, stride=1, padding=0, bias=False), # 71 x 71 x 192
            nn.MaxPool2d(3, stride=2, padding=0), # 35 x 35 x 192
        )
        self.branch_0 = Conv2d(192, 96, 1, stride=1, padding=0, bias=False)
        self.branch_1 = nn.Sequential(
            Conv2d(192, 48, 1, stride=1, padding=0, bias=False),
            Conv2d(48, 64, 5, stride=1, padding=2, bias=False),
        )
        self.branch_2 = nn.Sequential(
            Conv2d(192, 64, 1, stride=1, padding=0, bias=False),
            Conv2d(64, 96, 3, stride=1, padding=1, bias=False),
            Conv2d(96, 96, 3, stride=1, padding=1, bias=False),
        )
        self.branch_3 = nn.Sequential(
            nn.AvgPool2d(3, stride=1, padding=1, count_include_pad=False),
            Conv2d(192, 64, 1, stride=1, padding=0, bias=False)
        )
    def forward(self, x):
        x = self.features(x)
        x0 = self.branch_0(x)
        x1 = self.branch_1(x)
        x2 = self.branch_2(x)
        x3 = self.branch_3(x)
        return torch.cat((x0, x1, x2, x3), dim=1)
```

```
class Inception_ResNet_A(nn.Module):
    def __init__(self, in_channels, scale=1.0):
        super(Inception_ResNet_A, self).__init__()
        self.scale = scale
        self.branch_0 = Conv2d(in_channels, 32, 1, stride=1, padding=0, bias=False)
        self.branch_1 = nn.Sequential(
            Conv2d(in_channels, 32, 1, stride=1, padding=0, bias=False),
            Conv2d(32, 32, 3, stride=1, padding=1, bias=False)
        )
        self.branch_2 = nn.Sequential(
            Conv2d(in_channels, 32, 1, stride=1, padding=0, bias=False),
            Conv2d(32, 48, 3, stride=1, padding=1, bias=False),
            Conv2d(48, 64, 3, stride=1, padding=1, bias=False)
        )
        self.conv = nn.Conv2d(128, 320, 1, stride=1, padding=0, bias=True)
        self.relu = nn.ReLU(inplace=True)
    def forward(self, x):
        x0 = self.branch_0(x)
        x1 = self.branch_1(x)
        x2 = self.branch_2(x)
        x_res = torch.cat((x0, x1, x2), dim=1)
        x_res = self.conv(x_res)
        return self.relu(x + self.scale * x_res)
```



# Inception + Resnetv2

```
class Inception_ResNet_B(nn.Module):
    def __init__(self, in_channels, scale=1.0):
        super(Inception_ResNet_B, self).__init__()
        self.scale = scale
        self.branch_0 = Conv2d(in_channels, 192, 1, stride=1, padding=0, bias=False)
        self.branch_1 = nn.Sequential(
            Conv2d(in_channels, 128, 1, stride=1, padding=0, bias=False),
            Conv2d(128, 160, (1, 7), stride=1, padding=(0, 3), bias=False),
            Conv2d(160, 192, (7, 1), stride=1, padding=(3, 0), bias=False)
        )
        self.conv = nn.Conv2d(384, 1088, 1, stride=1, padding=0, bias=True)
        self.relu = nn.ReLU(inplace=True)
    def forward(self, x):
        x0 = self.branch_0(x)
        x1 = self.branch_1(x)
        x_res = torch.cat((x0, x1), dim=1)
        x_res = self.conv(x_res)
        return self.relu(x + self.scale * x_res)
```

```
class Inception_ResNet_C(nn.Module):
    def __init__(self, in_channels, scale=1.0, activation=True):
        super(Inception_ResNet_C, self).__init__()
        self.scale = scale
        self.activation = activation
        self.branch_0 = Conv2d(in_channels, 192, 1, stride=1, padding=0, bias=False)
        self.branch_1 = nn.Sequential(
            Conv2d(in_channels, 192, 1, stride=1, padding=0, bias=False),
            Conv2d(192, 224, (1, 3), stride=1, padding=(0, 1), bias=False),
            Conv2d(224, 256, (3, 1), stride=1, padding=(1, 0), bias=False)
        )
        self.conv = nn.Conv2d(448, 2080, 1, stride=1, padding=0, bias=True)
        self.relu = nn.ReLU(inplace=True)
    def forward(self, x):
        x0 = self.branch_0(x)
        x1 = self.branch_1(x)
        x_res = torch.cat((x0, x1), dim=1)
        x_res = self.conv(x_res)
        if self.activation:
            return self.relu(x + self.scale * x_res)
        return x + self.scale * x_res
```

```
class Inception_ResNetv2(nn.Module):
    def __init__(self, in_channels=3, classes=1000, k=256, l=256, m=384, n=384):
        super(Inception_ResNetv2, self).__init__()
        blocks = []
        blocks.append(Stem(in_channels))
        for i in range(10):
            blocks.append(Inception_ResNet_A(320, 0.17))
        blocks.append(Reduction_A(320, k, 1, m, n))
        for i in range(20):
            blocks.append(Inception_ResNet_B(1088, 0.10))
        blocks.append(Reduciton_B(1088))
        for i in range(9):
            blocks.append(Inception_ResNet_C(2080, 0.20))
        blocks.append(Inception_ResNet_C(2080, activation=False))
        self.features = nn.Sequential(*blocks)
        self.conv = Conv2d(2080, 1536, 1, stride=1, padding=0, bias=False)
        self.global_average_pooling = nn.AdaptiveAvgPool2d((1, 1))
        self.linear = nn.Linear(1536, classes)
    def forward(self, x):
        x = self.features(x)
        x = self.conv(x)
        x = self.global_average_pooling(x)
        x = x.view(x.size(0), -1)
        x = self.linear(x)
        return x
```


# Inception + Resnetv2

```
print(f'Accuracy on test set: {100 * correct / total:.2f}%')
```

```
⇒ Epoch [1/10], Loss: 0.3566, Accuracy: 87.61%  
Validation Accuracy: 86.36%  
Epoch [2/10], Loss: 0.2733, Accuracy: 90.17%  
Validation Accuracy: 93.69%  
Epoch [3/10], Loss: 0.2183, Accuracy: 91.38%  
Validation Accuracy: 94.58%  
Epoch [4/10], Loss: 0.1996, Accuracy: 92.16%  
Validation Accuracy: 88.09%  
Epoch [5/10], Loss: 0.2276, Accuracy: 91.74%  
Validation Accuracy: 87.87%  
Epoch [6/10], Loss: 0.2607, Accuracy: 89.30%  
Validation Accuracy: 91.02%  
Epoch [7/10], Loss: 0.1959, Accuracy: 92.40%  
Validation Accuracy: 96.80%  
Epoch [8/10], Loss: 0.1505, Accuracy: 94.41%  
Validation Accuracy: 96.53%  
Epoch [9/10], Loss: 0.1468, Accuracy: 94.60%  
Validation Accuracy: 97.60%  
Epoch [10/10], Loss: 0.1387, Accuracy: 94.82%  
Validation Accuracy: 97.16%  
Accuracy on test set: 96.62%
```

Epoch: 10

Accuracy: 96.62%



Predicción: adenocarcinoma con 99.06% confidence

# RESUMEN

Arquitectura	Nro de Épocas	Precisión Train	Precisión Test	Entorno de entrenamiento
Inceptionv1	10	97.69 %	97.24%	Local
ResNet 50	10	98%	98%	Google Colab
Inception + Resnetv2	10	94.82%	96.62%	Google Colab

**GPU** Colab T4 16GB

**GPU** local: GTX 1060 6GB

**Tiempo** Google Colab: 2h


**Tiempo** Local: 5h




# Inceptionv1

# Inception + Resnetv2


# ResNet-50



Predicción: adenocarcinoma con 97.00% confidence



Predicción: adenocarcinoma con 99.06% confidence



Predicción: adenocarcinoma con 98.48% confidence



**UNIVERSIDAD  
NACIONAL  
DE INGENIERÍA**

**Facultad de Ingeniería  
Industrial y de  
Sistemas**

**SECC. A**

# **CURSO REDES NEURONALES Y APRENDIZAJE PROFUNDO**

**GRACIAS**