


Computer Vision Nanodegree

Project: Image Captioning

The Microsoft **C**ommon **O**bjects in **C**ontext (MS COCO) dataset is a large-scale dataset for scene understanding. The dataset is commonly used to train and benchmark object detection, segmentation, and captioning algorithms.



You can read more about the dataset on the [website](#) or in the [research paper](#).

In this notebook, you will explore this dataset, in preparation for the project.

Step 1: Initialize the COCO API

We begin by initializing the [COCO API](#) that you will use to obtain the data.

```
In [1]: import os
import sys
sys.path.append('/opt/cocoapi/PythonAPI')
from pycocotools.coco import COCO

# initialize COCO API for instance annotations
dataDir = 'opt/cocoapi'
dataType = 'val2014'
instances_annFile = os.path.join(dataDir, 'annotations/instances_{}.json'.format(dataType))
coco = COCO(instances_annFile)

# initialize COCO API for caption annotations
captions_annFile = os.path.join(dataDir, 'annotations/captions_{}.json'.format(dataType))
coco_caps = COCO(captions_annFile)

# get image ids
ids = list(coco.anns.keys())
```

```
loading annotations into memory...
Done (t=2.41s)
creating index...
index created!
loading annotations into memory...
Done (t=0.15s)
creating index...
index created!
```

```
In [2]: list(coco.anns.values())[0]
```

```
Out[2]: {'segmentation': [[239.97,
    260.24,
    222.04,
    270.49,
    199.84,
    253.41,
    213.5,
    227.79,
    259.62,
    200.46,
    274.13,
    202.17,
    277.55,
    210.71,
    249.37,
    253.41,
    237.41,
    264.51,
    242.54,
    261.95,
    228.87,
    271.34]],
    'area': 2765.1486500000005,
    'iscrowd': 0,
    'image_id': 558840,
    'bbox': [199.84, 200.46, 77.71, 70.88],
    'category_id': 58,
    'id': 156}
```

Step 2: Plot a Sample Image

Next, we plot a random image from the dataset, along with its five corresponding captions. Each time you run the code cell below, a different image is selected.

In the project, you will use this dataset to train your own model to generate captions from images!

```
In [3]: import numpy as np
import skimage.io as io
import matplotlib.pyplot as plt
%matplotlib inline

# pick a random image and obtain the corresponding URL
ann_id = np.random.choice(ids)
img_id = coco.anns[ann_id]['image_id']
img = coco.loadImgs(img_id)[0]
url = img['coco_url']

# print URL and visualize corresponding image
print(url)
```

```
I = io.imread(url)
plt.axis('off')
plt.imshow(I)
plt.show()

# Load and display captions
annIds = coco_caps.getAnnIds(imgIds=img['id']);
anns = coco_caps.loadAnns(annIds)
coco_caps.showAnns(anns)
```

http://images.cocodataset.org/val2014/COCO_val2014_000000161962.jpg



A picture of a building and some grass.

A photo of a courtyard with a building in the background.

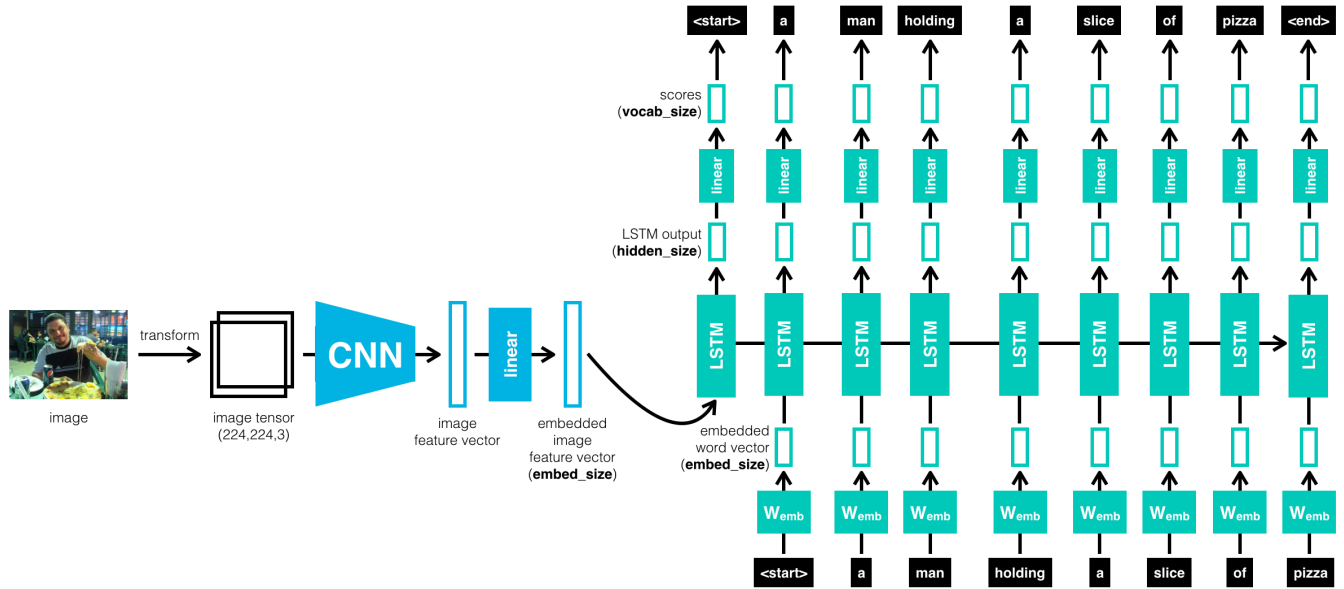
People are on the lawn enjoying the park situated in front of a tourist spot.

A castle-like building is in the background while the foreground is a green grass lawn, part of which has been mowed.

A large public park stationed in front of a building.

Step 3: What's to Come!

In this project, you will use the dataset of image-caption pairs to train a CNN-RNN model to automatically generate images from captions. You'll learn more about how to design the architecture in the next notebook in the sequence (**1_Preliminaries.ipynb**).



In []: